jax-experiments writeups Milestone M4

M4 — Parallel scans and compiled decoding (N4, O4)

A writeup in the jax-experiments series, built on 2026-10-30 from `docs/m4-scans-and-decode.md`. Its data, scripts and the work order (`PLAN.md`) are in the same repository.

Contents

N4 — Parallel scans: linear recurrences against the GRU and the transformer **1**
Question . . . 1
The idea . . . 2
The math . . . 2
What we measure and why . . . 3
What I built . . . 3
Results . . . 4
What surprised me . . . 7
Open threads . . . 7

O4 — KV-cached decoding, compiled **7**
Question . . . 8
The idea . . . 8
The math . . . 8
What we measure and why . . . 9
What I built . . . 9
Results . . . 9
What surprised me . . . 11
Open threads . . . 11

At a glance

Status: done (N4, O4) ·
Hardware: N4: Colab T4 (calibration, LRU sweeps, copy-LR check) and Kaggle T4 ×2 (Mamba-2 sweeps and seeds, copy task); O4: Colab T4; JAX 0.11.1; CPU (4 cores) for the equivalence tests ·
Data: N4: `data/n4-summary.json` and `data/n4/` (one JSON and log per job); O4: `data/o4-summary.json` and `data/o4/`; ml-ladder's references from its `docs/writeups/data/c-gpu.json`, `c-gpu-lr.json`, `c-gpu-lr2-256.json` and `c-gpu-lr3-gru-*.json` ·
Reproduce: N4: `uv run python tools/kaggle/push_r2.py <preset> --push` for `n4-calibrate`, `n4-lru`, `n4-mamba2`, `n4-mamba2-b`, `n4-mamba2-s0`, `n4-copy` and `n4-copy-mamba2`, then `uv run python experiments/n4_scans/analyze.py`; O4: preset `o4` (the command-buffer runs set `XLA_FLAGS`, see `docs/data/o4/o4-cmdbuf-*.log`), then `uv run python experiments/o4_decode/analyze.py`. CPU tests: `uv run pytest tests/test_n4_ssm.py tests/test_o4_decode.py`

N4 — Parallel scans: linear recurrences against the GRU and the transformer

Question

In ml-ladder's Part C the GRU was the best 4M model at T=256 and T=1024, but it walks the sequence one step at a time: cuDNN's GRU ran at 86k tokens/s at T=1024 against the 4M transformer's 136k (both measured here, in JAX, on one T4). Its state is fixed (two layers of 592 numbers), and fixed-state recurrences cannot copy a long sequence: R3's LSTM (two layers of 512) fails the copy task from n=128.

A *linear* recurrence keeps the fixed-size state but drops the nonlinearity between steps, and that lets it train in parallel. Three questions:

- **Speed and memory.** Does a parallel scan give transformer-like (or better) training speed with the RNN's flat memory profile?
- **Quality.** At about 4M parameters and ml-ladder's protocol, where does it land between the GRU and the transformer in bits/byte, and does its loss keep falling with position like the transformer's, or flatten like the GRU's?
- **Copying.** Does it fail the copy task, as a fixed-state model should?

The idea

A linear recurrence is a prefix product. The GRU's update $h_t = f(h_{t-1}, x_t)$ is nonlinear in h_{t-1}, so step t needs step $t-1$ first. A diagonal linear recurrence

$$h_t = a_t \odot h_{t-1} + b_t,$$

with a_t and b_t computed from the input alone, is different. Write each step as the pair (a_t, b_t), a map $h \mapsto a_t h + b_t$. Composing two such maps gives another one:

$$(a_1, b_1) \bullet (a_2, b_2) = (a_1 a_2,\ a_2 b_1 + b_2).$$

This operator is associative, so all T prefixes can be computed by a parallel scan in $O(\log T)$ depth and $O(T)$ work (`lax.associative_scan`). The nonlinearity moves out of the recurrence, into the layers between recurrences.

Three models, each about 4M parameters, sharing the byte-level shell of ml-ladder's language models (embedding, residual blocks, LayerNorm, linear head):

- **LRU** (Orvieto et al. 2023). A complex diagonal recurrence with a fixed transition $\lambda = \exp(-e^{\nu} + ie^{\theta})$, $|\lambda|$ initialised in $[0.9, 0.999]$. The input is normalised by $\sqrt{1-|\lambda|^2}$. Each block is an LRU mixer and a GELU MLP. The transition does not depend on the input, so the LRU decides *how long* to remember per channel, but not *what*.
- **Mamba-1 (S6)** (Gu & Dao 2023). The transition depends on the input: $a_t = \exp(\Delta_t A)$ and $b_t = \Delta_t B_t x_t$, with Δ_t, B_t, C_t computed from x_t ("selection"). The state is [channels × 16] per position. A causal depthwise convolution precedes it and a SiLU gate follows.
- **Mamba-2 (SSD)** (Dao & Gu 2024). The same selection, but with a *scalar* decay per head. That turns the scan into matrix products in chunks of 64 positions: an attention-like masked product inside each chunk, and a short recurrence across chunks.

Why Mamba-2 is here. Mamba-1 without its fused CUDA kernel materialises a [batch, T, 512, 16] state for the scan. The parallel scan then stores $O(\log T)$ copies of it for the backward pass. It ran, but at 13.7k tokens/s and 2 GiB with remat, and 12.4 GiB of temp memory without. That is a statement about the missing kernel, not about selective SSMs. Mamba-2's chunked form uses the matmul units that a T4 has instead, so it is the fair selective-SSM arm for the quality and copy comparisons.

The math

Cost per training step, for T positions and state width N per channel:

	work	sequential depth	activations stored
GRU (`lax.scan` / cuDNN)	$O(T)$	T	$O(T)$
linear scan, sequential	$O(T)$	T	$O(T)$
linear scan, associative	$O(T)$	$O(\log T)$	$O(T)$ (a few times more)
SSD, chunk Q	$O(TQ + TN)$	T/Q (across chunks)	$O(TQ)$ per head
attention	$O(T^2)$	1	$O(T^2)$ without a fused kernel

At a fixed 16,384 tokens per step (ml-ladder's protocol), the batch shrinks as T grows. A recurrence's cost per token is then independent of T; a sequential one's depth grows with T; attention's cost per token grows linearly in T.

Why a fixed state cannot copy. In the copy task, the model reads n random symbols from a 32-symbol vocabulary and must reproduce them. At the separator it must carry $n \log_2 32 = 5n$ bits in its state. The LRU's state is 256 complex numbers per layer, but each is a leaky average of the past, and storing a sequence in it needs distinct, slowly decaying channels per position. Nothing stops a linear model from learning a delay line for small n; the question is where the state runs out.

What we measure and why

- **Equivalence (CPU).** The associative scan equals a Python loop, for real and complex inputs; scan modes agree in value and gradient; chunked SSD equals the step-by-step recurrence (with padding, one chunk, and ragged lengths); all models are causal.
- **Speed and memory (T4).** Training tokens/s and XLA's compiled temp memory at T = 256, 1024 and 4096, fp16 matmuls, 16,384 tokens per step. Arms: the 4M transformer (fused attention), the GRU (`lax.scan`), LRU (associative and sequential), Mamba-1 (associative and sequential, remat) and Mamba-2 (chunked and sequential).
- **Quality (T4).** enwik8 bits/byte on ml-ladder's protocol (3,000 steps of 16,384 bytes, AdamW β=(0.9, 0.95), clipping at 1.0, 5% warmup then cosine to 10%; 262,144 validation bytes), at T=256 and T=1024. Weight decay is 0.1 for the new arms and the transformer, and 0 for the GRU (ml-ladder's settings). Each new arm gets its own LR sweep on seed 0 over {1e-3, 2e-3, 4e-3, 8e-3}, then seeds 0–2 at its best LR. The baselines are ml-ladder's own runs at their tuned LRs: GRU 2e-3 at both lengths; transformer 2e-3 at T=1024 and 8e-3 at T=256.

- **Loss by position.** Bits at each position in the window, median over seeds. Its summary is the *context gain*: the mean bits at positions $[T/8, T/4)$ minus those in the last quarter. It is what the second half of the window buys.
- **Copying.** R3's task: n = 32, 128, 512, 1,500 steps, second-half accuracy on 512 held-out sequences, seeds 0–2, each model at 2e-3 as in R3. For Mamba-2, which failed it, an LR check at n=32 (seed 0, 2.5e-4 to 8e-3) and one run of 6,000 steps.

What I built

- **`jaxexp.ssm`** — the three models as pure functions over dicts of arrays.
  - `linear_scan(a, b, mode)` runs the recurrence as `associative` (`lax.associative_scan`) or `sequential` (`lax.scan`).
  - `ssd(x, dta, b, c, chunk)` is Mamba-2's chunked form; `ssd_sequential` is the reference recurrence. `_segsum` builds the in-chunk decay matrix as a masked cumulative sum (a difference of two cumulative sums loses precision), with $-\infty$ above the diagonal so `exp` gives exact zeros. A test checks that no NaN from those entries reaches the gradients.
  - `apply(params, ids, kind, scan=..., compute_dtype=..., remat=...)`. `compute_dtype` applies to the dense projections only; recurrences, gates, LayerNorms and the residual stream stay in fp32.
- **Runners.** `experiments/r2_enwik8/run.py` and `experiments/n1_vmap_sweep/run.py` take `--model lru|mamba|mamba2` and `--scan`, so the LR sweeps run as vmapped groups (N1) on ml-ladder's batches for each seed, and the copy task runs on R3's code unchanged. `experiments/n4_scans/run.py` measures throughput and memory against T; `analyze.py` builds the tables and figure.
- **Tests.** `tests/test_n4_ssm.py`, 13 tests (above).

Results

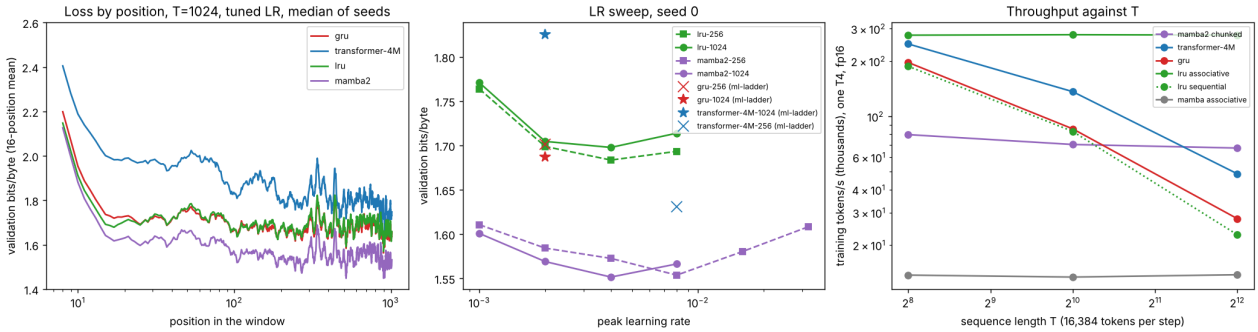


Figure 1. N4: loss by position at T=1024, LR sweeps, throughput against T

Left: validation bits/byte at each position of the 1,024-byte window, each model at its tuned LR, median of seeds 0–2 (16-position running mean). The GRU's and the transformer's curves are ml-ladder's own runs. Middle: seed 0's LR sweeps (lines) and ml-ladder's tuned baselines (markers). Right: training throughput on one T4 at a fixed 16,384 tokens per step.

Speed and memory (Colab T4, fp16 matmuls, 16,384 tokens per step). Training tokens/s, in thousands, and XLA's compiled temp memory for one step:

arm	T=256	T=1024	T=4096	temp memory (GiB)
transformer-4M (fused attention)	248	136	49	1.38 / 2.87 / 9.26
GRU (`lax.scan`)	196	86	28	1.01
LRU, associative scan	**276**	**278**	**277**	1.19
LRU, sequential scan	187	83	23	1.71
Mamba-2, chunked (SSD)	80	70	67	3.39 / 3.34 / 3.33
Mamba-1, associative (remat)	13.7	13.4	13.8	2.08 / 1.96 / 2.02
Mamba-2 or Mamba-1, sequential	out of memory			38.1 / 23.6

- **The parallel scan removes the dependence on T.** The associative LRU runs at 277k tokens/s at every length; the same recurrence run sequentially falls from 187k to 23k, the GRU's curve. At T=4096 it is 12× its sequential self, 10× the GRU and 5.7× the transformer.
- **Memory stays flat too:** 1.19 GiB at every T, against the transformer's 2.87 → 9.26 GiB as its T×T scores grow.
- **Mamba-2's chunked form is flat but slower:** 67–80k tokens/s, a quarter of the LRU, and a third (T=256) to a half (T=1024) of the transformer; it passes the transformer only at T=4096. Its step is a chain of einsums over chunk-by-chunk decay matrices and per-head states, 3.3 GiB of temp memory.
- **The sequential selective scans do not fit:** a `lax.scan` over Mamba-2's [heads, 64, 64] state stores it at every step for the backward pass, 38 GiB.
- **vmapped copies at T=1024** (temp GiB for K = 1, 2, 3 without remat): LRU 1.18 / 2.40 / 3.59, Mamba-2 3.69 / 6.63 / 9.95, Mamba-1 12.4 / 15.2 / 37.9. So the LRU sweeps ran as one program of 4 copies and Mamba-2's as groups of 2.

enwik8 (3,000 steps, ml-ladder's protocol). Seed 0's LR curves (bits/byte):

	1e-3	2e-3	4e-3	8e-3	1.6e-2	3.2e-2
LRU, T=1024	1.771	1.705	**1.698**	1.714		
LRU, T=256	1.764	1.699	**1.684**	1.693		
Mamba-2, T=1024	1.601	1.569	**1.552**	1.566		
Mamba-2, T=256	1.611	1.584	1.573	**1.554**	1.580	1.608

Every best LR is interior. Seeds 0–2 at the best LR, median (min–max), with ml-ladder's tuned baselines:

	T=256	T=1024
transformer-4M (ml-ladder; 8e-3 / 2e-3)	1.631 (1.624–1.634)	1.825 (1.823–1.828)
GRU (ml-ladder; 2e-3)	1.701 (1.700–1.701)	1.687 (1.687–1.692)
LRU (4e-3)	1.686 (1.684–1.688)	1.698 (1.693–1.699)
Mamba-2 (8e-3 / 4e-3)	**1.578 (1.554–1.581)**	**1.567 (1.552–1.571)**

- **The LRU is a GRU that trains in parallel.** It lands within 0.015 bits/byte of the GRU at both lengths (0.015 better at T=256, 0.011 worse at T=1024), at 3.2× the GRU's speed at T=1024.
- **Mamba-2 is the best 4M model here, by a wide margin:** 0.12 bits/byte better than the GRU at T=1024 and 0.05 better than the tuned transformer at T=256. At T=1024 it beats the transformer by 0.26.
- **Seed 0 is the best Mamba-2 run at both lengths, by about 0.02.** It is also the seed the LRs were picked on, so it was run again, alone, at both best LRs: 1.5595 at T=1024 (first run 1.5515) and 1.5513 at T=256 (1.5537). It reproduces within 0.008 and stays ahead of seeds 1–2 (1.567–1.581), so seed 0 is a better draw of initial weights and batches, not a lucky run. The medians above use the first runs.

Loss by position. The gap to the GRU at the same positions (bits/byte, median curves; both models are scored on the same validation windows, so the text's own difficulty at each position cancels):

T=1024	0–8	8–64	64–256	256–1024	trend
transformer-4M	+0.160	+0.255	+0.179	+0.118	−0.075
LRU	−0.036	−0.006	+0.006	+0.010	−0.000
Mamba-2	−0.056	−0.099	−0.120	−0.128	−0.004

Trend: the gap in the last quarter minus the gap at positions [T/8, T/4); negative means gaining on the GRU as the context grows.

- **The LRU behaves like the GRU at every position.** Its gap stays within 0.04 everywhere and does not trend: past the first ~64 bytes, neither uses more context.
- **Mamba-2 is ahead from the first bytes and pulls away up to a few hundred bytes,** then holds a constant lead: its gap grows from −0.06 to −0.12 by position 256 and only to −0.13 after. So it uses more context than the GRU, but not the whole window.
- **Only the transformer keeps gaining** (trend −0.075 at T=1024, −0.065 at T=256). At T=1024 it starts so far behind that it is still 0.12 worse than the GRU in the last three quarters; at T=256 it is level with the GRU over the first 8 bytes and 0.09 ahead by positions 64–256.
- **Raw per-position curves are noisy.** Bin to bin, the GRU's and the LRU's curves move together by ±0.03 with the text, so a model's own "context gain" (early bits minus late bits) is noise at the ±0.02 level; only the paired gap is informative.

Copying (R3's task, 1,500 steps, second-half accuracy; chance 0.031).

n	LSTM (R3)	GRU	LRU	Mamba-2 (2e-3)	transformer-4M (R3)
32	0.893 (0.864–0.910)	0.886 (0.851–0.917)	**1.000** (all)	0.065 (0.056–0.996)	1.000
128	0.031	0.202 (0.201–0.208)	**1.000** (all)	0.031	1.000
512	0.031	0.040 (0.034–0.040)	0.148 (0.132–0.159)	0.031	1.000 (0.981–1.000)

- **The LRU copies 128 symbols perfectly**, where both gated RNNs fail; at n=512 it gets 15%, five times chance. Its state is 256 complex numbers per layer; 128 symbols are $5 \times 128 = 640$ bits, 512 symbols 2,560.
- **Mamba-2 fails even at n=32 on two seeds of three.** This is optimisation, not capacity. A sweep on seed 0 at n=32 (second-half accuracy): 2.5e-4: 0.153; 5e-4: 0.137; 1e-3: 0.130; 2e-3: 0.065; 4e-3: 0.029; 8e-3: 0.030. Lower LRs plateau around 0.14, higher ones stay at chance. Its loss falls slowly through all 1,500 steps (3.63 → 3.21 nats at 5e-4), so the same run was given 6,000 steps, four times the protocol's budget. **It then copies: 0.979** (loss 1.95 nats at the end; the floor is 1.71). Mamba-2 can copy 32 symbols; it needs most of 6,000 steps to get there, where the LRU is done by step 375.
- **The transformer remains the only model that copies 512.**

What surprised me

- **The LRU learns to copy almost at once.** At n=32 its training loss is at the task's floor by step 375, the first one logged: 1.71 nats, which is $\ln 32 = 3.47$ on the unpredictable first half and about 0 on the copy. The GRU is still at 1.88 after 1,500 steps, and Mamba-2 at 3.39. The LRU is causal (`tests/test_n4_ssm.py` changes one input and checks every earlier output is unchanged) and both scan modes agree, so this is not a leak. A fixed, input-independent linear recurrence can hold a delay line: a channel with $|\lambda| \approx 1$ and a phase that rotates by a fixed angle per step keeps each past symbol at a recoverable angle. Copying is the task such a memory is built for, and the GRU's gates have to learn it.
- **The best language model here is the slowest to learn to copy.** Mamba-2 leads on enwik8 and trails on the copy task, the opposite ranking to the LRU's. A plausible reading, not tested: its decays are chosen from the input, which text rewards, while holding 32 random symbols needs decays held near 1 for exactly the copy's duration, a setting the LRU starts with and Mamba-2 must find.
- **Mamba-1 without its kernel is a memory problem first.** Its associative scan materialises the [B, T, 512, 16] state at every level of the scan tree. Remat brings 12.4 GiB to 2 GiB, but the step is then 20× slower than the LRU's. Mamba-2 exists for this reason: the scalar decay per head turns the scan into matmuls.
- **Run-to-run noise on the GPU is up to 0.014 bits/byte for Mamba-2.** The same T=256 run at 1.6e-2 gave 1.567 on Colab and 1.580 on Kaggle; the seed-0 repeats moved by 0.008 and 0.002. So the 0.02 gap between seed 0 and seeds 1–2 is real, but a single-run difference under ~0.015 is not evidence.

Open threads

- **Equal compute, not equal steps.** Mamba-2 runs at a third to a half of the transformer's tokens/s at T ≤ 1024, so at equal GPU-hours the transformer would get two to three times the steps. Its lead at equal steps may not survive that comparison; it was not run.
- **Why the LRU stops at a few hundred bytes of context.** Its |λ| is initialised in [0.9, 0.999], so its longest memory starts at about 1,000 steps; whether training moves it, and whether r_max = 0.9999 helps T=1024, was not tried.
- **Copying at n=128 and 512 with a longer budget for Mamba-2.** At n=32 it needed 6,000 steps; whether it reaches n=128, where the LRU is perfect, was not run.
- **A fused kernel for Mamba-1's selective scan** (Pallas, on a TPU; Kaggle's GPUs cannot run Pallas, see CLAUDE.md). The question it answers: is the 20× gap to the LRU the missing kernel, or the [channels × state] state?

O4 — KV-cached decoding, compiled

Question

ml-ladder's C3 and C5 explain why a decoder caches keys and values: without a cache every new token recomputes the whole prefix, $O(T^2)$ work per token; with one it reads the cache once, $O(T)$. C5 adds that decoding is then memory-bound: per token, the GPU must read every weight and the whole cache, and arithmetic is trivial next to that. C5's open thread asked for the check: the bytes a step reads, divided by the card's bandwidth, against measured tokens/s at several context lengths.

In JAX the question has a second half. A cached decode step is a small jitted function called once per token. How much does it cost to call it from Python, and does putting the whole generation inside one compiled program (`lax.fori_loop`) reach the bandwidth floor? That is what the candle fork's static decode did with preallocated buffers replayed in a CUDA graph (davidmiheev/candle#3 and #4).

The idea

Static shapes. XLA compiles one program per input shape, so a cache that grows by one token per step would recompile every step. Instead the cache is preallocated for C positions, the step writes the new key and value at position p with `lax.dynamic_update_slice` and attends over all C slots under a mask that hides positions after p. Every step has the same shape, so one compiled program serves the whole generation. The price is that attention reads all C slots, used or not.

Three decoders, all greedy so that they must emit the same tokens:

- **recompute** — no cache: a full forward pass over a $[B, C]$ buffer per token, reading the logits at the current position;
- **cached** — prefill the prompt in one parallel pass, then one jitted step per token, called from a Python loop (the cache buffer is donated, so XLA updates it in place);
- **loop** — prefill and all steps inside one jitted `lax.fori_loop`.

The math

Bytes per cached step at batch B, context C, L layers of width d, with fp16 weights and cache:

$$\text{bytes} = 2N + 2 \cdot 2LBCd,$$

the N weights read once, plus the whole static cache, keys and values. At the T4's 320 GB/s this gives the *bandwidth floor* on a step's time. For the 25M RoPE model ($L = 8$, $d = 512$): weights 51 MB; the cache is 16.8 MB at $C = 1024$, $B = 1$, and 537 MB at $C = 4096$, $B = 8$, where it is ten times the weights.

Recompute does a full forward pass over $B \times C$ tokens per generated token: $2NBC$ FLOPs for the matmuls plus $O(LBC^2d)$ for attention. It is compute-bound and grows at least linearly in C, quadratically at long contexts.

Launch overhead. A GPU kernel launch costs a few microseconds. A decode step of the 25M model compiles to 157 kernels (the 4M: 105). At batch 1 the whole step should take 0.2 ms by bandwidth, so 157 launches at ~4 µs each (0.6 ms) would dominate it. A CUDA graph records the launches once and replays them as one; XLA's equivalent is its command buffers.

What we measure and why

- **Equivalence (CPU).** The functional forward pass equals the NNX model; each cached step's logits equal the full forward pass at that position (to 10^{-5} in fp32); the three decoders emit identical tokens.
- **Decode cost (T4).** For R2's 25M RoPE transformer at C = 256 to 4096 and the 4M (learned positions, so C ≤ 1024), batch 1 and 8, plus batch 32 at C = 1024: ms per step after compilation (median of 3), compile time, prefill time, XLA's compiled temp memory, and XLA's own estimate of bytes accessed per step. A step decodes one token per sequence; 128 tokens are generated from a prompt of $C - 128$.
- **Kernel count.** Top-level instructions of the compiled step's optimised HLO.
- **Command buffers.** The cached and loop arms with XLA's command buffers at their default, off, and enabled for every command type including `WHILE`.

Weights are random: a step's cost does not depend on their values.

What I built

- **`jaxexp.decode`.** `weights(model)` reads a `CharTransformer` into plain arrays, matmul weights in fp16 once, up front. `forward` is the model's forward pass over those arrays, also returning every layer's keys and values. `prefill`, `step`, `empty_cache`, and the three generators `make_recompute`, `make_cached`, `make_loop`. Numerics follow `transformer.TORCH`; scores and softmax are fp32, as in R2.
- **`experiments/o4_decode/run.py`** and `analyze.py`; Kaggle preset `o4`.
- **Tests.** `tests/test_o4_decode.py`, 8 tests.

Results

Milliseconds per decode step (one new token for each of B sequences) against the static context length C, on one T4 (Colab), fp16. Dashed: the bandwidth floor, the bytes a cached step must read divided by 320 GB/s.

The 25M RoPE transformer, batch 1 (ms per token):

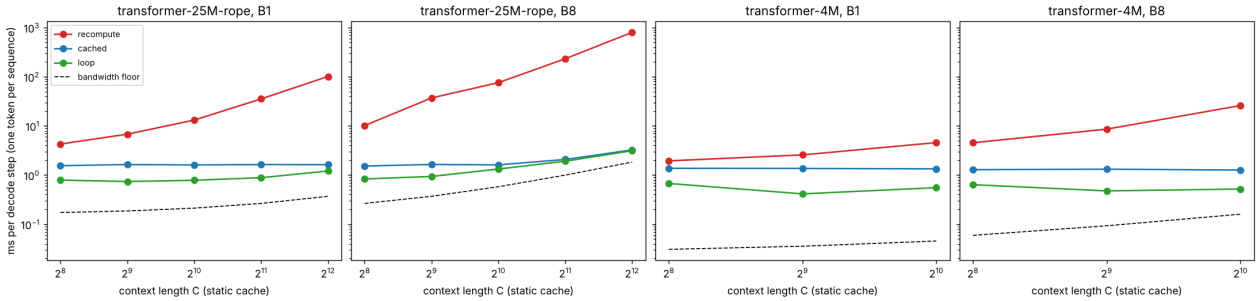


Figure 2. O4: ms per decode step against context length

C	recompute	cached (Python loop)	loop (`fori_loop`)	bandwidth floor	prefill (ms)
256	4.28	1.55	0.79	0.17	3.1
512	6.79	1.64	0.74	0.19	5.5
1024	13.13	1.61	0.79	0.21	12.1
2048	35.44	1.64	0.88	0.27	33.2
4096	101.10	1.63	1.21	0.37	95.5

Batch 8 (ms per step of 8 tokens):

C	recompute	cached	loop	floor
256	10.18	1.53	0.83	0.27
1024	75.82	1.61	1.33	0.58
4096	800.54	3.25	3.13	1.84

and batch 32 at C=1024 (the same 588 MB per step as batch 8 at C=4096): cached 3.30, loop 3.21, floor 1.84.

- **The cache turns quadratic into flat.** Recompute grows from 4.3 to 101 ms per token at batch 1 as C goes from 256 to 4096 (24× for 16× the context), and to 801 ms at batch 8. The cached step stays at 1.5–1.6 ms. At C=4096 the cache is 62× faster at batch 1 and 246× at batch 8.
- **At batch 1 nothing is bandwidth-bound.** The floor is 0.17–0.37 ms. The cached step called from Python costs 1.55–1.64 ms whatever C is; the compiled loop halves that (0.74–0.88 ms up to C=2048) and is still 3–5× the floor. At C=4096 the loop rises to 1.21 ms as the cache read starts to count.
- **Only large steps reach the floor's neighbourhood.** When a step must read about 300 MB or more (batch 8 at C ≥ 2048, batch 32 at C=1024), the cached and loop arms converge (3.25 / 3.13 ms for 588 MB at batch 8; 3.30 / 3.21 at batch 32) at 178–188 GB/s, 56–59% of the T4's 320 GB/s.
- **The bandwidth arithmetic matches XLA's.** XLA's own bytes-accessed estimate for a cached step is within 1–9% of the model's weights + whole cache at batch 1 for the 25M (1.02–1.09×). For the 4M it is 1.7× (fixed costs dominate a 10 MB step), and at batch 8 with short contexts it is lower than the model (0.52× at C=256); I did not trace which fusions it undercounts.
- **The 4M model is launch-bound everywhere:** cached 1.27–1.37 ms, loop 0.42–0.68 ms, against a floor of 0.03–0.16 ms.

- **Every decoder emits the same tokens** in 14 of the 17 configurations. In the other three, fp16 near-ties: recompute disagrees with the cache on 3.5% and 3.1% of tokens (25M, B=8, C=256 and 2048) and 2.7% (4M, B=8, C=1024). The cached and loop decoders, which run the same step, always agree.
- **Compile time and memory.** The step compiles in 2.9–5.6 s, the full loop in 3.9–7.8 s (it includes the prefill), recompute in 1.7–4.6 s. With its cache donated, the cached step needs 0.3 MiB of temp memory at batch 1: the update happens in place. The loop's temp memory is its prefill's: 890 MiB at C=4096, batch 1, and 6.8 GiB at batch 8.

Where the loop's time goes. The step compiles to 157 kernels for the 25M model (105 for the 4M). The loop's time above the bandwidth floor, divided by the kernel count, is 3.5–3.9 µs per kernel for the 25M at C ≤ 2048, and 3.6–6.1 µs for the 4M: a fixed cost per kernel.

XLA's command buffers (CUDA graphs), batch 1, ms per token:

	cached, 25M C=1024	loop, 25M C=1024	cached, 25M C=4096	loop, 25M C=4096	cached, 4M	loop, 4M
off	2.05	0.75	2.08	1.20	1.65	0.63
default	1.71	0.79	1.70	1.26	1.39	0.67
all, incl. `WHILE`	1.57	0.89	1.62	1.30	1.44	0.53

- **Command buffers help the step called from Python** (2.05 → 1.57 ms): the step's launches are recorded once and replayed.
- **They do not help the loop.** Off, default and all land within 0.14 ms of each other, in no consistent order, and that is the run-to-run spread: the 4M's default loop measured 0.555 and 0.669 ms in two separate runs. So the loop's remaining ~0.55 ms is not host-side launch overhead that a graph can remove; it is the device's fixed cost per small kernel. Only fewer kernels would lower it.

What surprised me

- **Python costs more than the model at batch 1.** Dispatching a jitted step with a pytree of 117 leaves (100 weight arrays and 17 LayerNorm epsilons) from Python takes about 0.8 ms; the step itself, inside the loop, about the same. The candle fork's CUDA graph attacked the same cost; here a `fori_loop` gets most of the way without one.
- **The CUDA-graph analogue does not finish the job.** XLA's command buffers, even with the while loop inside them, leave the loop where it was. On a T4, 157 small kernels at ~3.6 µs each are a floor of their own, three times the bandwidth floor.
- **The static cache reads empty slots and it rarely matters.** Attention reads all C positions every step whether used or not. At batch 1 that is invisible (launch-bound); it starts to show only at C=4096 (loop 0.79 → 1.21 ms).

Open threads

- **Fewer kernels.** The loop's floor is per-kernel cost. A fused decode step (Pallas, on a TPU or an Ampere GPU; the T4 cannot run Pallas) or fusing the QKV, rotary and attention of one layer would test how far below 0.75 ms batch 1 can go.

- **A bucketed cache** (attend over the next power of two above the current length via `lax.switch`) would stop reading empty slots; at batch 8 and C=4096 it could halve the bytes for the first half of a generation.
- **PyTorch's numbers on the same card.** ml-ladder's decode measurements (C5's four real decoders) used other models; the same 25M model in PyTorch with and without CUDA graphs would separate framework from hardware.